

Containerisation with Docker

Aurélien Coet
March 2020



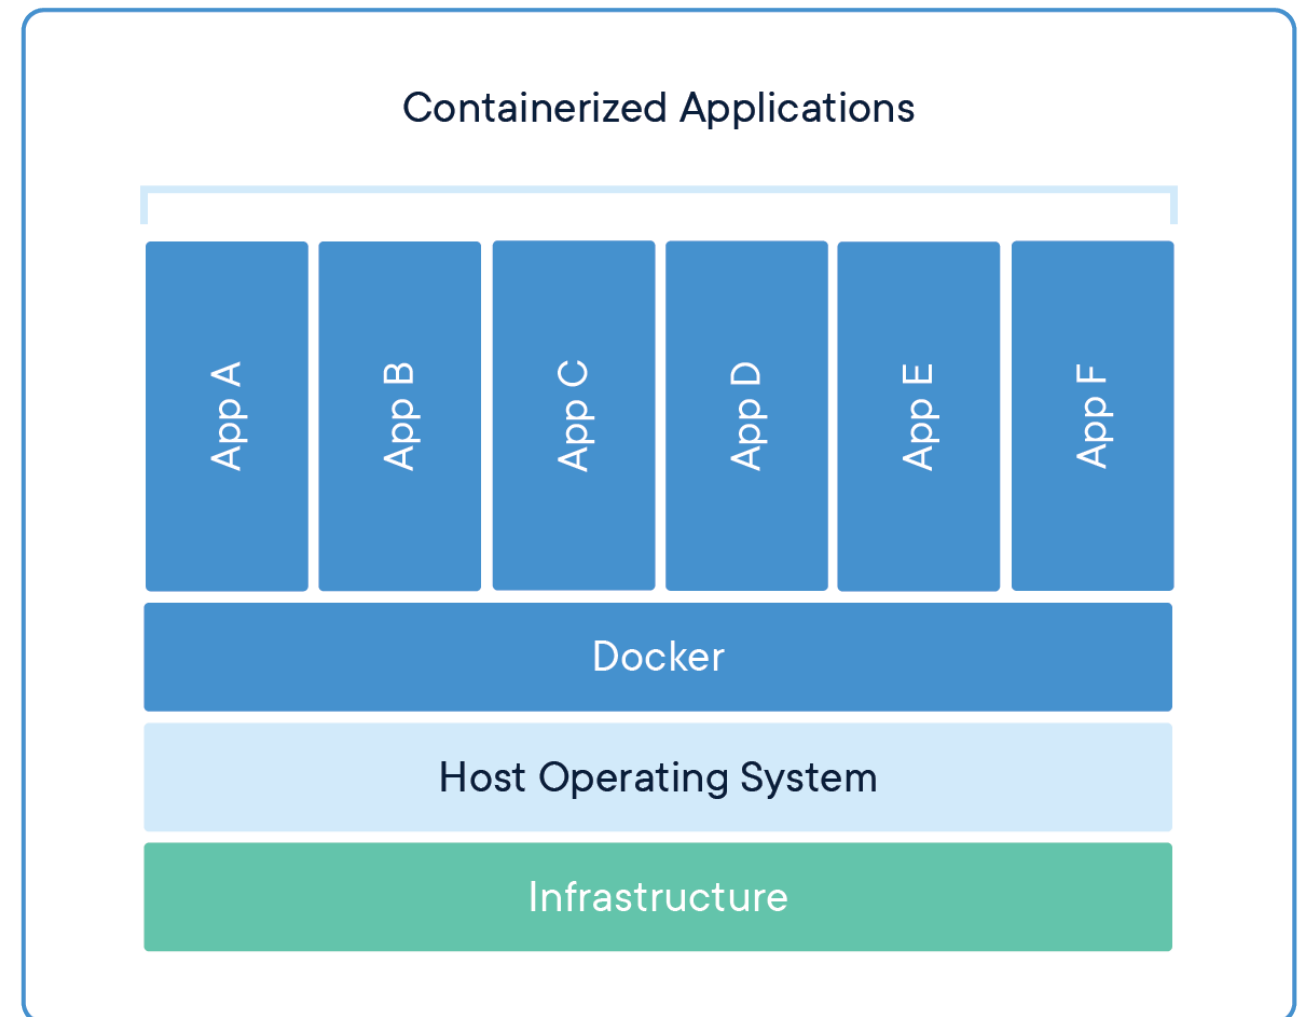
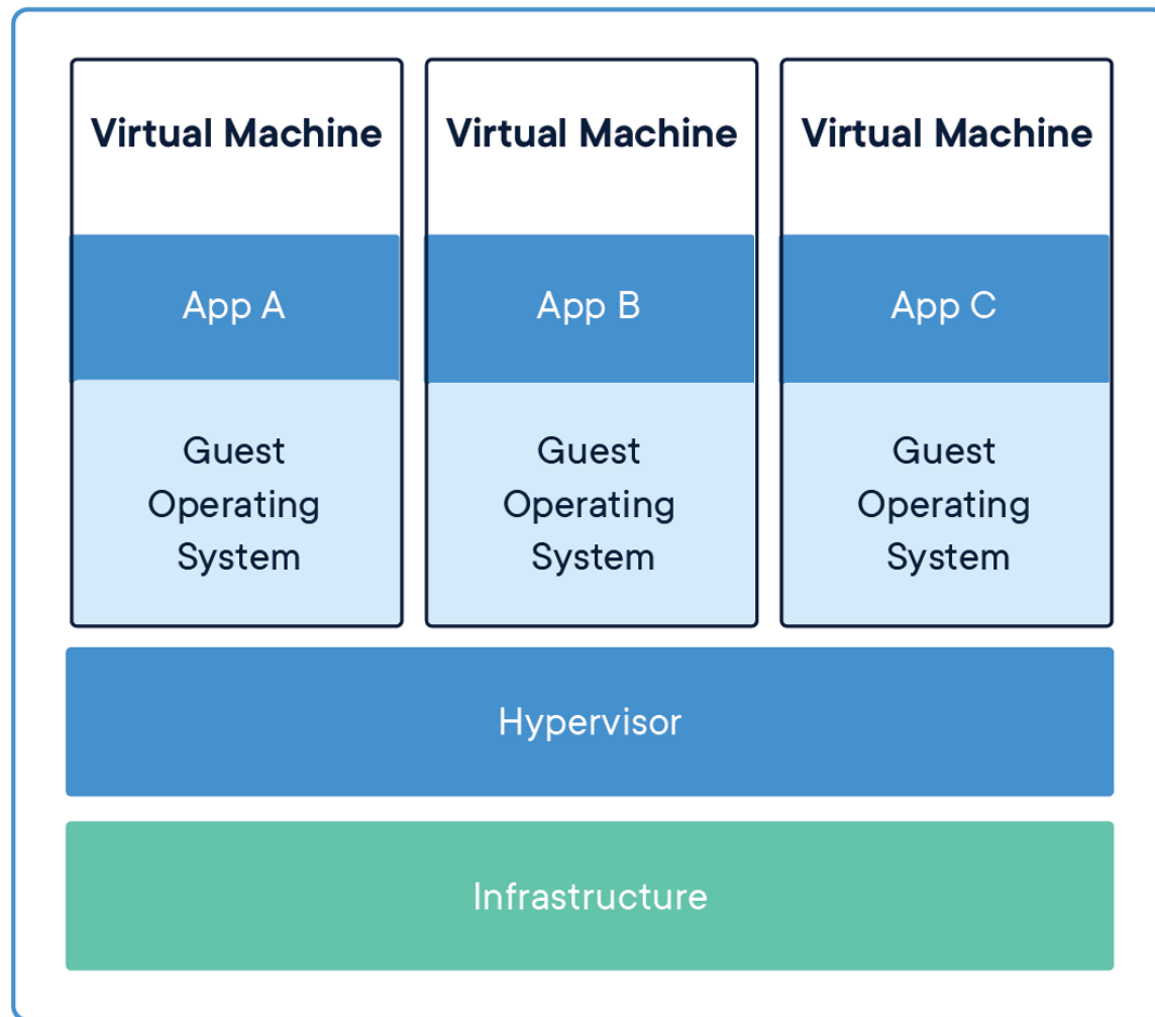
**UNIVERSITÉ
DE GENÈVE**
FACULTÉ DES SCIENCES

What is *Docker* ?

What is *Docker* ?

- A tool designed to make software deployment easier by bundling it in *containers*
- Containers are software units that can be directly run on a machine, without additional installation or configuration steps
- They package together configuration files, libraries and programs in isolated environments

Virtual Machines vs Docker Containers



Docker Basics

Docker Concepts

- Docker containers are built from *images*
- Images define what a container will contain and execute when it is spawned
- They are like *templates* for containers: an image is unique, and multiple containers can be spawned from the same image

Docker Images

- In the *microservices* example:

unige/regulatory-service	latest	5859668ecfb1	12 seconds ago	778MB
unige/valuation-service	latest	93516633b7b3	48 seconds ago	814MB
unige/instrument-service	latest	b1bded92050c	About a minute ago	814MB
unige/counterparty-service	latest	1789c8543673	2 minutes ago	780MB
unige/api-gateway	latest	b355613b0bbd	32 hours ago	371MB

are docker images containing the microservices and their dependencies (JVM, libraries, etc...)

- *myCounterpartyService* is a container started from the counterparty-service image

Docker Containers

- Docker containers are spawned with the command:

```
$ docker run -p x:y --name=z img:tag
```

where:

- **x:y** is a mapping from OS port x to container port y
- **z** is the name given to the container
- **img** is the name of the image from which the container must be spawned
- **tag** is the tag (version) of the image

Docker Containers

- When **docker run** is called on an image, the *docker daemon* first searches for it locally on the user's machine
- If it doesn't find any image with the input name, it looks for it in a *docker registry*
- By default, the registry in which docker looks up for images is **DockerHub**

Docker Containers

- Running docker containers are listed with the command:

```
$ docker ps
```

- All the containers on a machine can be listed with:

```
$ docker ps -a
```

Docker Containers

- Docker containers are stopped with the command:

```
$ docker kill id
```

where:

- *id* is the name of the container to stop (or its auto-generated identifier if no name was provided when it was started)
- All running docker containers can be stopped with:

```
$ docker kill $(docker ps -q)
```

Docker Containers

- Docker containers are deleted with the command:

```
$ docker rm id
```

where:

- *id* is the name of the container to stop (or its auto-generated identifier if no name was provided when it was started)
- All existing docker containers can be stopped with:

```
$ docker rm $(docker ps -a -q)
```

Dockerfiles

- Docker images are defined with *Dockerfiles*

```
FROM nginx
COPY ./dist /usr/share/nginx/html
COPY nginx.default /etc/nginx/conf.d/default.conf
COPY httpasswd /etc/nginx/httpasswd
```

<https://docs.docker.com/engine/reference/builder/>

Dockerfiles - Build stages

- Dockerfiles are split into *build stages* (or *layers*) that define a set of actions to incrementally build an image
- Build stages start from a *base image* defining the environment of the stage
- Each stage can use a different base image
- The instruction:

```
FROM nginx AS stage-name
```

- Defines *nginx* as the stage's base image
- Names the stage *stage-name*

Dockerfiles - Build stages

```
FROM node:latest AS build-stage  
...  
FROM nginx:latest AS deploy-stage  
...
```

} 1st layer

} 2nd layer

More on build stages:

<https://docs.docker.com/develop/develop-images/multistage-build/>

Dockerfiles - Workdirs

- Build stages operate on *working directories*

- The instruction:

```
WORKDIR somedir
```

- Defines *somedir* as the stage's workdir
- Creates the directory in the image if it doesn't exist

Dockerfiles - Copying files

- The instruction:

```
COPY ./localdir ./somedir
```

- Copies the files in *localdir* on the user's machine to the build stage's *somedir* working directory

Dockerfiles - Running commands

- The instruction:

RUN command

- Runs the command *command* inside the image

.dockerignore

- ***.dockerignore* files work similarly to *.gitignore* files in git**
- **They indicate to the docker daemon that it must ignore some files on the user's machine when building an image**
- **This prevents leaking sensitive or useless data into docker images**

Building a docker image

- A docker image is built with the command:

```
$ docker build -t name:tag path
```

where:

- *name* is the name to give to the image
- *tag* is the tag for the image
- *path* is the path to a Dockerfile

Publishing a docker image

- A docker image is pushed to the docker registry with the command:

```
$ docker push name:img
```

where:

- *name* is the name of the image
- *tag* is the tag of the image

Listing docker images

- All docker image on the machine can be listed with the command:

```
$ docker image ls
```

Deleting a docker image

- A docker image is deleted with the command:

```
$ docker rmi -t name:tag
```

where:

- *name* is the name of the image
- *tag* is the tag for the image
- All existing docker images can be deleted with:

```
$ docker rmi $(docker image ls)
```
- An image cannot be deleted if a container using it is running !

Docker command line reference:

[**https://docs.docker.com/engine/reference/commandline/docker/**](https://docs.docker.com/engine/reference/commandline/docker/)

Dockerfile reference:

[**https://docs.docker.com/engine/reference/builder/**](https://docs.docker.com/engine/reference/builder/)